SPIRE — Tester Introduction

What you're testing: an early build of SPIRE, a system that turns a plain-English question about your Oracle Fusion pod into a real answer pulled live from Fusion.

What it is not: a chatbot, an assistant, a co-pilot. It is a question-answer tool with a deliberately narrow scope. It only answers questions about the Fusion modules and record types listed below. Out-of-scope questions get a polite refusal, not a guess.

How to access

- **URL:** *<your ngrok URL here — update per session>*

Open the URL in any browser. No login. The page shows a green **Connected** badge when it can reach the backend. If it's red, wait a moment and refresh.

The URL changes each time the tunnel is restarted. If a link stops working, ask for the new one.

Please do not share this URL further. It reaches a live demo pod with query limits.

How to ask a question

Type your question in plain English. Press Submit.

- **Good:** "Show suppliers in the US with no invoices in 6 months."
- **Good:** "Top 5 suppliers by spend this year."
- **Good:** "Show PO 2144 and its lines."

- **Not useful:** "What can you do?" — SPIRE does not converse.
- **Not useful:** "Write me a summary of the supplier contract." — not a data question, out of scope.

What SPIRE supports

Fusion modules

Four. Questions about other modules (GL, HCM, SCM, projects, etc.) will be refused.

Module	What's in scope
Supplier	Supplier master, addresses, sites, contacts, business classifications
Procurement	Purchase orders, PO lines, PO schedules, PO distributions, buyers
Payables	Accounts-payable invoices
Receivables	AR invoices, AR invoice lines, AR invoice tax lines, AR invoice flexfields, customer accounts, customer transactions, credit memos, receipts, receipt applications

Common things you can ask for

Suppliers

- List suppliers by status (active/inactive)
- Find suppliers by country, name fragment, business relationship
- Suppliers with a site in a specific business unit
- Suppliers with a specific business classification
- Suppliers with open purchase orders
- Suppliers with no recent invoices
- Suppliers with an invoice over a specific amount

Purchase orders

- Show a specific PO by number, including its lines, schedules, and distributions
- POs by supplier, by status (open, closed, cancelled, etc.)
- POs created before or after a date
- Top suppliers by committed PO value
- PO status breakdown

Payables invoices

- Invoices for a supplier (by name or number)
- Invoices in a date range, in an amount range
- Invoices raised against a specific PO

Receivables

- AR invoices by invoice number
- AR invoices with their lines and tax lines
- Customer accounts by customer name, with balance and activity
- Customer transactions for an account
- Credit memos (list, or per-account)
- Receipts (list, or per-account, filtered by customer or site)
- Receipt applications

Aggregations

- Count of suppliers / invoices / POs
- Top N suppliers by spend, by PO value
- Receipts grouped by method
- Customer accounts ranked by open exposure

What it will refuse

- Anything outside the four modules above
- Questions requiring a data source other than Oracle Fusion REST APIs
- Compound questions that need data the pods do not expose
- Vague questions without a clear entity or filter

The refusal message will tell you **why** and give you an example of what to ask instead. Please share those with us — they're often the signal we need.

What to look for while testing

Three things, in priority order.

1. Wrong answers. The pipeline cross-checks every ID in the answer against the source data, so an invented ID should be impossible. If you see one anyway, that's the most valuable bug you can report.

2. Empty answers that shouldn't be empty. "No suppliers found" when you know suppliers exist. The pipeline is honest about emptiness — but the filters might be too narrow.

3. Refusals on questions that should work. SPIRE tries to answer honestly rather than guess; a refusal on a question you think is in scope is useful feedback.

Known limits (not bugs)

- **Truncation.** Aggregate questions scan a page-capped slice of the pod. The answer says so plainly when it happened — look for "Showing N of M" in the answer text. Not a defect.
- **No history.** Each question is independent. The page does not remember previous questions.
- **No follow-ups.** "Show me the same but for the UK" will not work. Ask the full question each time.
- **Latency varies.** Some questions scan a lot of pod data and take 30–60 seconds. The UI shows a spinner.
- **The pod is a demo instance.** Some data may be incomplete or

synthetic. If something looks suspicious but is internally consistent, tell us; don't assume.

How to report feedback

For each item please include:

1. **The exact question you asked** (copy-paste)
2. **What you saw** (screenshot or copied answer text)
3. **What you expected** (one sentence)
4. **When** (roughly — the pod data changes)

Send to:

Short feedback is fine. "Asked X, got Y, expected Z" is enough.